

Data Management

Anthony Chau

UCI Center for Statistical Consulting

2021/01/12 (updated: 2021-01-27)

Data Management

- Now, we'll address the **manipulation problem**: how to select and change slices of our data
- The focus is on data frames but other data structures will be discussed

Data management is a broad topic, so I'll focus on a few common tasks

1. Select specific columns
2. Create a new column
3. Filter data given a condition
4. Rename column
5. Group data into subsets


Background

- R provides **subsetting** operators that allow us to select data in complex and useful ways
- **Subsetting** is the action of selecting specific pieces of our data
- How we subset data is dependent on the data type and data structure

Subsetting Operators : `[]`, `[[]]`, `$`

Vector mental model

- Recall our vector mental model



Subset a vector with the [operator

- General syntax: `v[...]`
- Within the brackets, we can provide a integer, character, or logical vector
- Supply an integer vector to select by index
- Supply an character vector to select by element name
- Supply an logical vector to select by condition

Subset a vector with integers

- Subset with positive or negative integers
- Use `c()` to subset with a vector of length > 1

```
x ← c(-1, 0, 2, 3)

# notice the use of c() inside the bracket for vector length > 1

# select the first element
x[1]
#> [1] -1

# select first and fourth element
x[c(1, 4)]
#> [1] -1 3

# exclude first and fourth element
x[c(-1, -4)]
#> [1] 0 2

# can't combine positive and negative indices
x[c(-1, 2)]
#> Error in x[c(-1, 2)]: only 0's may be mixed with negative subscripts
```

Subset a vector with element names

- Subset with element names

```
month_days ← c(31, 28, 31, 30, 31)
names(month_days) ← c("Jan", "Feb", "Mar", "Apr", "May")

# c() is not required for a length one subsetting vector
month_days["Feb"]
#> Feb
#> 28
month_days[c("Jan", "Apr")]
#> Jan Apr
#> 31 30
```

Aside: logical operators

- R has built-in **logical operators** (operators used to evaluate whether a condition is true or false)

Operator	Description
<	less than
<=	less than or equal to
>	greater than
>=	greater than or equal to
==	exactly equal to

Logical operators example

```
x ← c(-10, -1, 0, 2, 3)

# remember vector recycling: (-10, -1, 0, 2, 3) > (0, 0, 0, 0, 0)
# x > 0 returns a logical vector of the same length as z
x > 0
#> [1] FALSE FALSE FALSE TRUE TRUE

# select and return all elements greater than 0
x[x > 0]
#> [1] 2 3

# select and return all elements less than or equal to 0
x[x ≤ 0]
#> [1] -10 -1 0

# select and return all elements equal to -10
x[x == -10]
#> [1] -10

# select and return all elements greater than 5
x[x > 5]
#> numeric(0)
```

Aside: boolean operators

- R has built-in **boolean operators** (operators used to chain together multiple logical expressions)

Operator	Description
!x	NOT x
x y	x OR y
x & y	x AND y

! boolean operator

- ! reverses the logical value (TRUE becomes FALSE, FALSE becomes TRUE)

```
# assume that foods are either fruits or vegetables
x ← c("apple", "spinach", "broccoli", "blueberry", "carrot")

fruit ← c(TRUE, FALSE, TRUE, TRUE, FALSE)
fruit
#> [1] TRUE FALSE TRUE TRUE FALSE

# return fruits
x[fruit]
#> [1] "apple"      "broccoli" "blueberry"

# reverse each logical value in fruit
vegetable ← !fruit
vegetable
#> [1] FALSE TRUE FALSE FALSE TRUE

# return vegetables
x[vegetable]
#> [1] "spinach" "carrot"
```

| and & boolean operator

- ! evaluates to TRUE if at least one logical expression is true
- & evaluates to TRUE if and only if all logical expressions are true
- Recommended to use parentheses to separate logical expressions

```
# first expression is TRUE; second expression is FALSE
(1 < 5)
#> [1] TRUE
(1 < -2)
#> [1] FALSE
# at least one expression is TRUE, so entire expression is TRUE
(1 < 5) | (1 < -2)
#> [1] TRUE
# not all expressions are TRUE, so entire expression is FALSE
(1 < 5) & (1 < -2)
#> [1] FALSE
```

Subset a vector with conditions

```
mascots ← c("Peter", "Tommy", "King Triton",  
            "Josephine", "Oski", "King Triton")  
names(mascots) ← c("UCI", "USC", "UCSD", "UCLA", "UCB", "UCSD")  
uc_campus ← c(TRUE, FALSE, TRUE, TRUE, TRUE, TRUE)  
  
# select elements that equal "Peter" or elements that equal "King Triton"  
mascots[mascots = "Peter" | mascots = "King Triton"]  
#>           UCI           UCSD           UCSD  
#> "Peter" "King Triton" "King Triton"  
  
# select non-UC campuses  
!uc_campus  
#> [1] FALSE  TRUE FALSE FALSE FALSE FALSE  
mascots[!uc_campus]  
#>      USC  
#> "Tommy"  
  
# select elements with element name "UCSD"  
names(mascots) = "UCSD"  
#> [1] FALSE FALSE  TRUE FALSE FALSE  TRUE  
mascots[names(mascots) = "UCSD"]  
#>           UCSD           UCSD  
#> "King Triton" "King Triton"
```

6 ways to subset a vector

Method	Behavior	Example	Result	Notes
Positive Integers	Select elements at the specified index	<code>x[c(1, 4)]</code> <code>x[c(1, 1)]</code>	<i>Return first and fourth element</i> <i>Return first element twice</i>	Duplicate indices return duplicate values Real numbers truncated to integers
Negative Integers	Exclude elements at the specified index	<code>x[c(-1, -4)]</code> <code>x[c(-2, 2)]</code>	<i>Exclude first and fourth element</i> <i>Error - not possible</i>	Can't mix positive and negative integer indices
Logical Vectors	Select elements when logical value is TRUE	<code>x[c(TRUE, FALSE, TRUE)]</code> <code>x[x > 0]</code>	<i>Return first and third element</i> <i>Return elements that are greater than 0</i>	
Nothing	Return the original vector	<code>x[]</code>	<i>Return the original vector</i>	Not that useful for vectors
Zero	Return a zero-length vector	<code>x[0]</code>	<i>Return empty numeric vector</i>	
Character Vectors	Select elements with matching names	<code>x[c("a", "c", "d")]</code>	<i>Return elements with element names:</i> <i>"a", "c", "d"</i>	Vector must have element names

Subset a list with [

- **Subsetting a list with [will always return a list**
- Just like vectors, you can supply a vector when using [

```
l <- list(letter = "a",
          number = 1,
          truthy = TRUE,
          ones_vector = c(1,1,1),
          my_list = list(1,2,3))

length(l)
#> [1] 5
names(l)
#> [1] "letter"      "number"      "truthy"      "ones_vector" "my_list"

# select the first element
l[1]
#> $letter
#> [1] "a"
is.list(l[1])
#> [1] TRUE

# select the element named "truthy"
l["truthy"]
#> $truthy
#> [1] TRUE
is.list(l["truthy"])
#> [1] TRUE
```

Subset a list with [

- All the ways to subset a vector carry through when subsetting a list with [

```
# vectors allowed
l[c("truthy", "number")]
#> $truthy
#> [1] TRUE
#>
#> $number
#> [1] 1
is.list(l[c("truthy", "number")])
#> [1] TRUE

# negative integers allowed
# exclude the second through fifth elements
l[c(-2, -3, -4, -5)]
#> $letter
#> [1] "a"
is.list(l[c(-2, -3, -4, -5)])
#> [1] TRUE
```


Subset a list with [[]]

- **Subsetting a list with [[]]** returns a single element in the list (the element *could* be a list)
- When using [[]], you can supply a single positive integer, a single element name, or a vector
- If you use a vector with [[]], you will subset recursively

```
l ← list(letter = "a", number = 1, truthy = TRUE, num_vector = c(1,2,3))
# single positive integer
l[[1]]
#> [1] "a"
is.list(l[[1]])
#> [1] FALSE
# single name
l[["truthy"]]
#> [1] TRUE
is.list(l[["truthy"]])
#> [1] FALSE

# recursive indexing: l[[c(4,3)]] = l[[4]][[3]]
l[[c(4, 3)]]
#> [1] 3
# no negative integers
l[[-2]]
#> Error in l[[-2]]: invalid negative subscript in get1index <real>
```

Subset a list with \$

- Subsetting a list with \$ is a shorthand for subsetting with [[
- `l$element_name = l[["element_name"]]`

```
l ← list(letter = "a", number = 1, truthy = TRUE, ones_vector = c(1,1,1))

l$letter
#> [1] "a"
is.list(l$letter)
#> [1] FALSE

l[["letter"]]
#> [1] "a"
is.list(l[["letter"]])
#> [1] FALSE
```

Subset a matrix with [

- Subsetting a matrix with [is similar to subsetting a vector with [
- Since a matrix is 2-dimensional, we select rows and columns with `m[row, column]`
- Then, we can provide a vector for each dimension to select specific rows and columns.

```
m ← matrix(1:16, nrow = 4, ncol = 4)
colnames(m) ← c("a", "b", "c", "d")
m
#>      a b  c  d
#> [1,] 1 5  9 13
#> [2,] 2 6 10 14
#> [3,] 3 7 11 15
#> [4,] 4 8 12 16

# first row, second column
m[1, 2]
#> b
#> 5

# first and third row; column a and column c
m[c(1, 3), c("a", "c")]
#>      a  c
#> [1,] 1  9
#> [2,] 3 11
```

Matrix subsetting shortcuts

- Syntax to **select all rows**, `m[, columns]`
- Syntax to **select all columns**, `m[rows,]`

```
m ← matrix(1:16, nrow = 4, ncol = 4)
colnames(m) ← c("a", "b", "c", "d")
m
#>      a b  c  d
#> [1,] 1 5  9 13
#> [2,] 2 6 10 14
#> [3,] 3 7 11 15
#> [4,] 4 8 12 16

# all rows; first and third column
m[, c(1, 3)]
#>      a  c
#> [1,] 1  9
#> [2,] 2 10
#> [3,] 3 11
#> [4,] 4 12

# first and second row; all columns
m[c(1, 4), ]
#>      a b  c  d
#> [1,] 1 5  9 13
#> [2,] 4 8 12 16
```

Single index subsetting

- Another way to subset a matrix is with a single vector
- Each element in a matrix is stored in column-major order

column major order:

start at top-left corner -> move down a column -> ...
start at top of adjacent column

```
m ← matrix(1:16, nrow = 4, ncol = 4)
colnames(m) ← c("a", "b", "c", "d")
m
#>      a b  c  d
#> [1,] 1 5  9 13
#> [2,] 2 6 10 14
#> [3,] 3 7 11 15
#> [4,] 4 8 12 16

# select first element and eleventh element
m[c(1,11)]
#> [1]  1 11
```

Subset a data frame

- Subsetting a data frame combines subsetting features from vectors, lists, and matrices
- Subsetting a data frame with a single index `df[]` is similar to list subsetting
- Subsetting a data frame with two indices `df[row, column]` is similar to matrix subsetting

Single index subsetting

- Subsetting a data frame with a single index selects the columns of a data frame

```
mascots <- data.frame(name = c("Peter Anteater", "Josephine Bruin",  
                               "King Triton", "Tommy Trojan"),  
                      age = c(56, 101, 60, 140),  
                      gpa = c(4, 3.9, 3.87, 3.7),  
                      residence = c("Irvine", "Los Angeles",  
                                   "San Diego", "Los Angeles"))
```

```
mascots[2]  
#>   age  
#> 1  56  
#> 2 101  
#> 3  60  
#> 4 140  
is.data.frame(mascots[2])  
#> [1] TRUE
```

Single index subsetting

- Single index subsetting for data frames is similar to vectors

```
# subset with double vector
mascots[c(1,3)]
#>      name  gpa
#> 1 Peter Anteater 4.00
#> 2 Josephine Bruin 3.90
#> 3   King Triton 3.87
#> 4   Tommy Trojan 3.70
is.data.frame(mascots[c(1,3)])
#> [1] TRUE
# subset with character vector
mascots[c("age", "gpa")]
#>   age  gpa
#> 1  56 4.00
#> 2 101 3.90
#> 3  60 3.87
#> 4 140 3.70
is.data.frame(mascots[c("age", "gpa")])
#> [1] TRUE
```


Double index subsetting

- Subsetting a data frame with two indices selects the rows and columns of a data frame

```
# select second and third row; all columns
mascots[c(2,3), ]
#>           name age  gpa  residence
#> 2 Josephine Bruin 101 3.90 Los Angeles
#> 3   King Triton  60 3.87  San Diego
is.data.frame(mascots[c(2,3), ])
#> [1] TRUE
# select all rows; name and residence columns
mascots[, c("name", "residence")]
#>           name  residence
#> 1 Peter Anteater    Irvine
#> 2 Josephine Bruin Los Angeles
#> 3   King Triton  San Diego
#> 4   Tommy Trojan Los Angeles
is.data.frame(mascots[, c("name", "residence")])
#> [1] TRUE
```

Double index subsetting

- Note that selecting only one column with double index subsetting does not return a data frame - this can be a source of downstream problems

```
# returns a data frame
mascots["age"]
#>   age
#> 1  56
#> 2 101
#> 3  60
#> 4 140
is.data.frame(mascots["age"])
#> [1] TRUE

# returns a vector
mascots[, "age"]
#> [1] 56 101 60 140
is.data.frame(mascots[, "age"])
#> [1] FALSE
```

Applications of subsetting

- Knowing the mechanics of subsetting opens up a variety of new operations we can perform on our data
 - We will focus on the following operations for data frames
1. Select specific columns
 2. Create a new column
 3. Filter data given a condition
 4. Rename a column
 5. Group data into subsets

Select specific columns

- Select columns with [, [[or \$

```
df <- data.frame(a = c(1:3), b = c(4:6), c = c(7:9), d = c(10:12))

df[, c("a", "d")]
#>   a  d
#> 1 1 10
#> 2 2 11
#> 3 3 12
df[c("b", "c")]
#>   b  c
#> 1 4  7
#> 2 5  8
#> 3 6  9
# use drop = FALSE to return a data frame
df[, c("b"), drop=FALSE]
#>   b
#> 1 4
#> 2 5
#> 3 6
df[["b"]]
#> [1] 4 5 6
df$b
#> [1] 4 5 6
```

Create a new column

- Create a new column by combining subsetting and assignment
- On the left hand side of the assignment operator, supply the name of your new column
- On the right hand side of the assignment operator, In general, apply some transformation on your existing columns to derive a new column

```
# raw data: weight ~ pounds, height ~ meters
df <- data.frame(height_m = c(1.7, 1.65, 1.9, 1.8, 1.73, 1.7),
  weight_lbs = c(151, 149, 187, 183, 175, 178),
  age = c(22, 20, 21, 19, 20, 19),
  sex = c("Male", "Male", "Female",
    "Male", "Male", "Female"),
  home_state = c("CA", "AZ", "TX", "CA", "FL", "NY"),
  school = c("USC", "UCI", "UCLA",
    "USC", "Chapman", "UCSD"))
```

```
df
#>   height_m weight_lbs age    sex home_state school
#> 1    1.70      151   22  Male        CA      USC
#> 2    1.65      149   20  Male        AZ      UCI
#> 3    1.90      187   21 Female        TX     UCLA
#> 4    1.80      183   19  Male        CA      USC
#> 5    1.73      175   20  Male        FL Chapman
#> 6    1.70      178   19 Female        NY     UCSD
```

Create a new column

- Note that after creating a new column, it is immediately available for use

```
# convert weight to kg
df[["weight_kg"]] <- df$weight_lbs / 2.205
df
#>   height_m weight_lbs age    sex home_state school weight_kg
#> 1    1.70      151  22  Male      CA      USC    68.48073
#> 2    1.65      149  20  Male      AZ      UCI    67.57370
#> 3    1.90      187  21 Female      TX     UCLA    84.80726
#> 4    1.80      183  19  Male      CA      USC    82.99320
#> 5    1.73      175  20  Male      FL Chapman  79.36508
#> 6    1.70      178  19 Female      NY     UCSD    80.72562
# compute bmi
df[["bmi"]] <- df$weight_kg / (df$height_m)^2
df
#>   height_m weight_lbs age    sex home_state school weight_kg    bmi
#> 1    1.70      151  22  Male      CA      USC    68.48073  23.69575
#> 2    1.65      149  20  Male      AZ      UCI    67.57370  24.82046
#> 3    1.90      187  21 Female      TX     UCLA    84.80726  23.49231
#> 4    1.80      183  19  Male      CA      USC    82.99320  25.61518
#> 5    1.73      175  20  Male      FL Chapman  79.36508  26.51779
#> 6    1.70      178  19 Female      NY     UCSD    80.72562  27.93274
```

Aside: conditional values with `ifelse()`

- Often, you want to create a new column given some condition in your existing columns
- The `ifelse()` function allows us to return a value given some condition
- The arguments to `ifelse()` are:
 - `test`: a logical expression that evaluates to TRUE or FALSE
 - `yes`: the value returned if test is TRUE
 - `no`: the value returned if test is FALSE

```
x ← c(-2, -5, 0, -2, 3, 4)

ifelse(test = x < 0 & x ≠ 0,
       yes = "Negative",
       no = "Positive")
#> [1] "Negative" "Negative" "Positive" "Negative" "Positive" "Positive"
```

Create a new column on a condition

- Let's create a column to indicate if a student is from out-of-state

```
# create an out-of-state column
df["out_of_state"] ←
  ifelse(test = df$home_state ≠ "CA", yes = "Yes", no = "No")

df
#>   height_m weight_lbs age    sex home_state school weight_kg    bmi out_of_state
#> 1    1.70      151    22  Male      CA      USC   68.48073 23.69575      No
#> 2    1.65      149    20  Male      AZ      UCI   67.57370 24.82046      No
#> 3    1.90      187    21 Female      TX     UCLA   84.80726 23.49231      No
#> 4    1.80      183    19  Male      CA      USC   82.99320 25.61518      No
#> 5    1.73      175    20  Male      FL Chapman 79.36508 26.51779      Yes
#> 6    1.70      178    19 Female      NY     UCSD   80.72562 27.93274      Yes
```


Filter data given a condition

- Filter data by combining logical expressions with subsetting
- Often, we want to see only the rows of our data that match a specific condition
- The general syntax to filter using subsetting operators:

```
df[logical_expression, selected_columns]
```

Filter data given a condition

- Let's view all the students who are older than 20

```
df[df$age > 20, ]  
#>   height_m weight_lbs age  sex home_state school weight_kg      bmi out_o  
#> 1      1.7      151  22  Male      CA      USC  68.48073 23.69575  
#> 3      1.9      187  21 Female     TX     UCLA  84.80726 23.49231
```

- Let's view all the male students

```
df[df$sex == "Male", ]  
#>   height_m weight_lbs age  sex home_state school weight_kg      bmi out_o  
#> 1      1.70      151  22  Male      CA      USC  68.48073 23.69575  
#> 2      1.65      149  20  Male     AZ     UCI  67.57370 24.82046  
#> 4      1.80      183  19  Male     CA      USC  82.99320 25.61518  
#> 5      1.73      175  20  Male     FL Chapman 79.36508 26.51779
```

subset()

- It turns out there is a built-in function `subset()` that allows you to select and filter data all in the same function

```
subset(x = df, subset = sex == "Male", select = c("home_state", "school"))
#>   home_state school
#> 1         CA    USC
#> 2         AZ    UCI
#> 4         CA    USC
#> 5         FL Chapman

# default to return all columns
subset(x = df, subset = age > 20)
#>   height_m weight_lbs age  sex home_state school weight_kg  bmi out_
#> 1      1.7      151  22  Male         CA    USC  68.48073 23.69575
#> 3      1.9      187  21 Female        TX   UCLA  84.80726 23.49231
```

Rename a column

- Rename columns by subsetting the column names of a data frame and assigning new names

```
df2 <- data.frame(EmaIl = c("chaua3@uci.edu",  
                           "peteranteater@uci.edu"),  
                  School.name = c("UCI", "UCI"))  
  
names(df2)  
#> [1] "EmaIl"      "School.name"  
  
# first, select the column names you want to change  
names(df2)[names(df2) == "EmaIl"]  
#> [1] "EmaIl"  
  
# then, assign the old name with a new name  
names(df2)[names(df2) == "EmaIl"] <- "email"  
  
names(df2)  
#> [1] "email"      "School.name"
```

Group data into subsets

- A common workflow is to group data and apply an operation (mean, sum) on each group
- Let's consider some sales data for a business. Each row is a transaction

```
sales <- data.frame(
  month = c("Jan", "Jan", "Jan", "Jan", "Feb", "Feb", "Feb", "Feb",
            "Mar", "Mar", "Mar", "Mar"),
  date = c("2020-01-05", "2020-01-24", "2020-01-14", "2020-01-13",
            "2020-02-01", "2020-02-04", "2020-02-07", "2020-02-14",
            "2020-03-02", "2020-03-14", "2020-03-13", "2020-03-21"),
  sales = c(13, 17, 21, 15, 12, 14, 17, 13, 21, 15, 14, 19))
# view only first four rows
head(sales, 4)
#>   month      date sales
#> 1  Jan 2020-01-05    13
#> 2  Jan 2020-01-24    17
#> 3  Jan 2020-01-14    21
#> 4  Jan 2020-01-13    15
```

Group data into subsets

- Group data and apply a summary function on each group with `aggregate()`
- Important arguments to `aggregate()`:
 - `x`: specify name of summary column and columns to apply summary function on
 - `by`: specify name and values of grouping column
 - `FUN`: specify summary function

```
# total sales for each month
aggregate(x = list(total_sales = sales$sales),
          by = list(month = sales$month),
          FUN = sum)
#>   month total_sales
#> 1   Feb           56
#> 2   Jan           66
#> 3   Mar           69

# average sales for each month
aggregate(x = list(average_sales = sales$sales),
          by = list(month = sales$month),
          FUN = mean)
#>   month average_sales
#> 1   Feb          14.00
#> 2   Jan          16.50
#> 3   Mar          17.25
```